

Martin Braun - Bachelor Thesis presentation



THURSDAY, 24 September 2015



Integration of JPA-conform ORM- Implementations in Hibernate Search

CONTACT

Martin Braun
E-Mail martinbraun123@aol.com

1. Introduction

- Abstraction/standardization is key in the software world (Java Enterprise)

- Database access is standardized
 - Object Relational Mappers (ORM)
 - Java Persistence API (JPA)
 - Hibernate ORM (Red Hat)
 - EclipseLink (Eclipse Foundation)
 - OpenJPA (Apache Foundation)

- Special features of different JPA implementations
 - e.g. Hibernate Search



THURSDAY, 24 SEPTEMBER 2015

CONTACT

Martin Braun
E-Mail martinbraun123@aol.com

1. Introduction

What is Hibernate Search?

- Lucene based fulltext search engine on top of Hibernate ORM
 - indexes and queries managed JPA objects
 - keeps index up-to-date

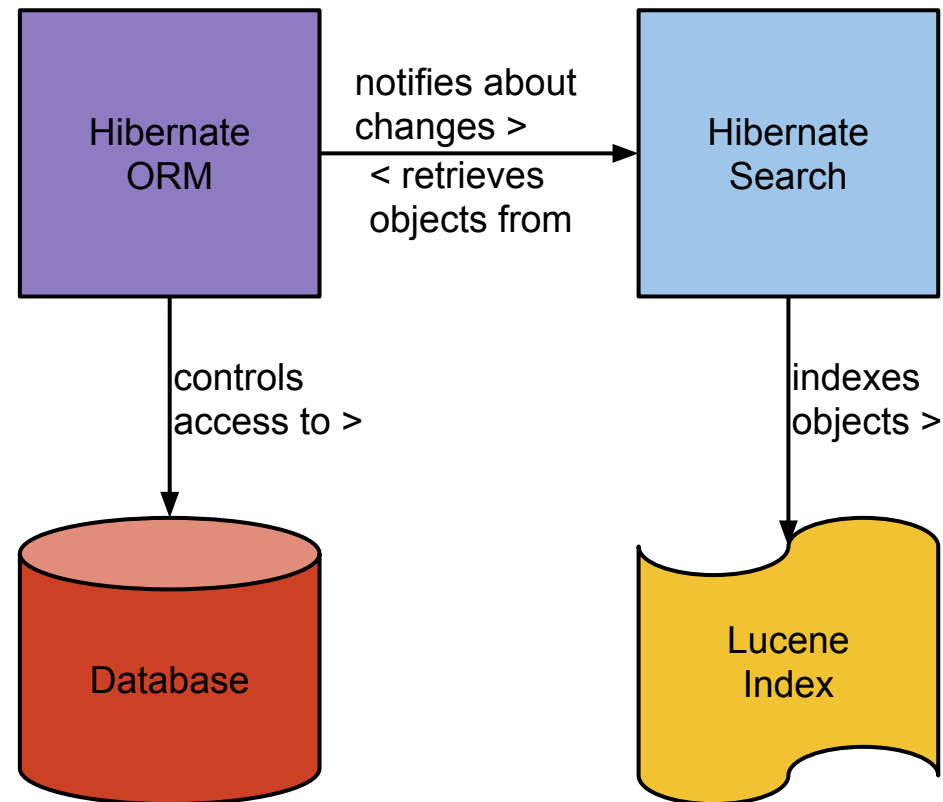
- Fulltext search in a regular RDBMS:

SELECT book.id, book.name FROM book WHERE book.name LIKE %hobbit%;

- Hibernate Search: more complex queries and index options
 - fuzzy queries *original hits original*
 - regular expression queries *[LI]ucene hits Lucene, lucene*
 - stemming (language specific) *worker -> work*
 - comprehensive synonym support *book -> essay, album, novel*

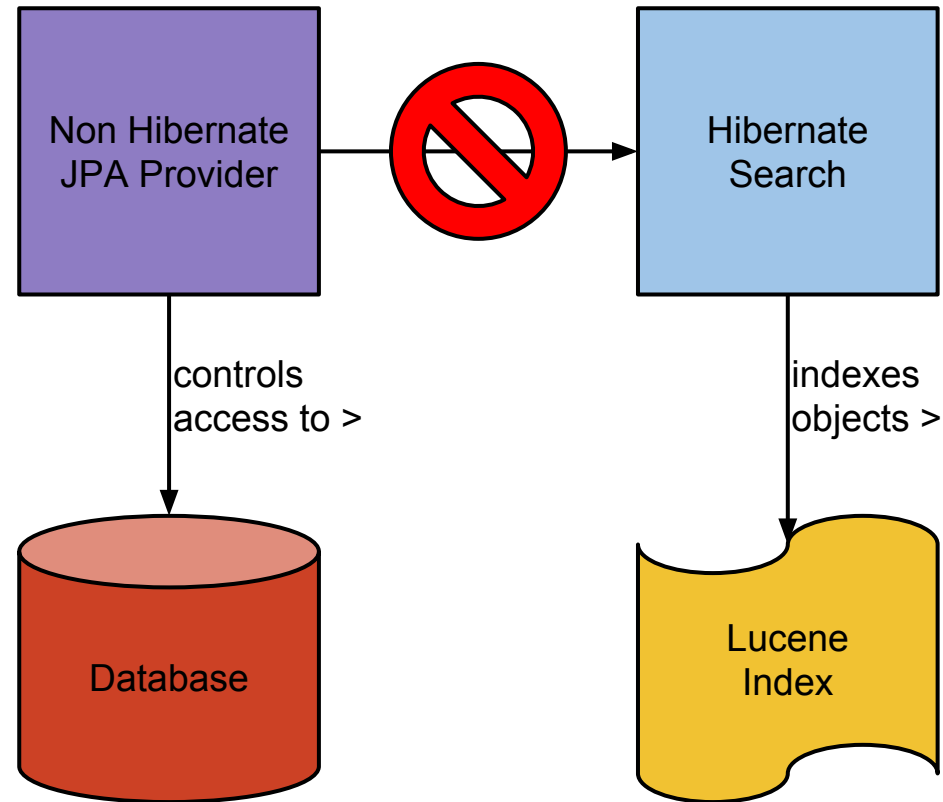
1. Introduction

Hibernate Search with Hibernate ORM:



1. Introduction

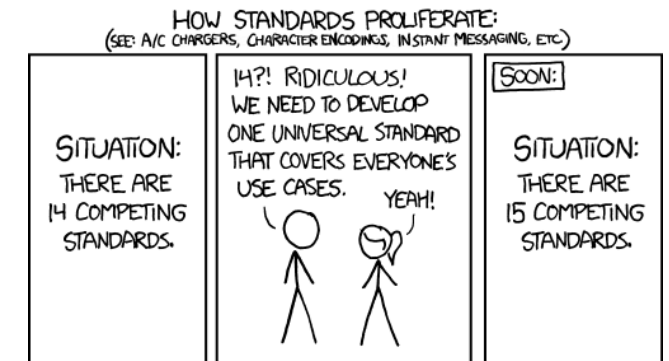
Hibernate Search with other JPA providers:



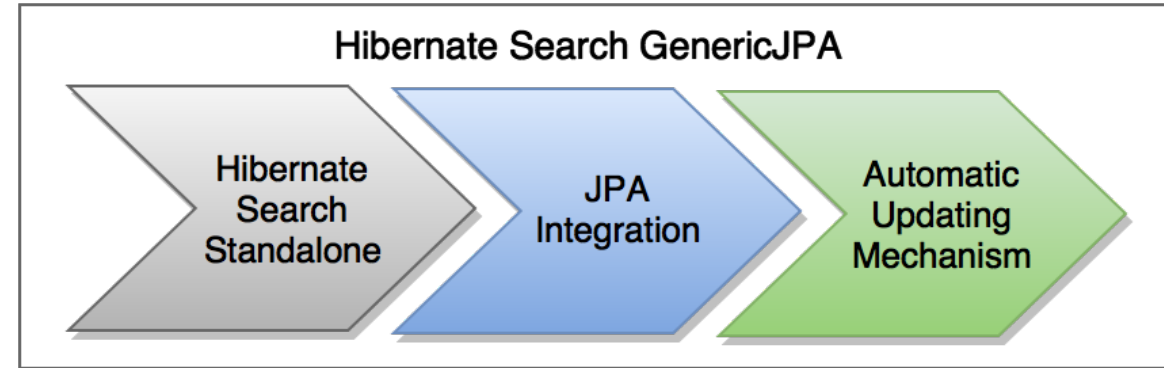
1. Introduction

Why a generic Hibernate Search?

- Other search solutions are not integrated with JPA:
 - native Lucene
 - ElasticSearch / Solr
- lack of integration results in:
 - manual conversion into index objects required
 - no automatic index updating
 - no JPA objects as return values in queries
- No new generic alternative needed:
 - existing JPA based interfaces in Hibernate Search
 - backend (hibernate-search-engine) is integration agnostic
- This approach is backed by the Hibernate team




2. Challenges



- **Hibernate Search Standalone**
 - abstraction layer on-top of the low-level API of Hibernate Search's engine
- **JPA integration**
 - integration of the standalone with generic JPA
 - re-use the interfaces from hibernate-search-orm
- **Automatic Updating Mechanism**



3. Standalone Version of Hibernate Search



```

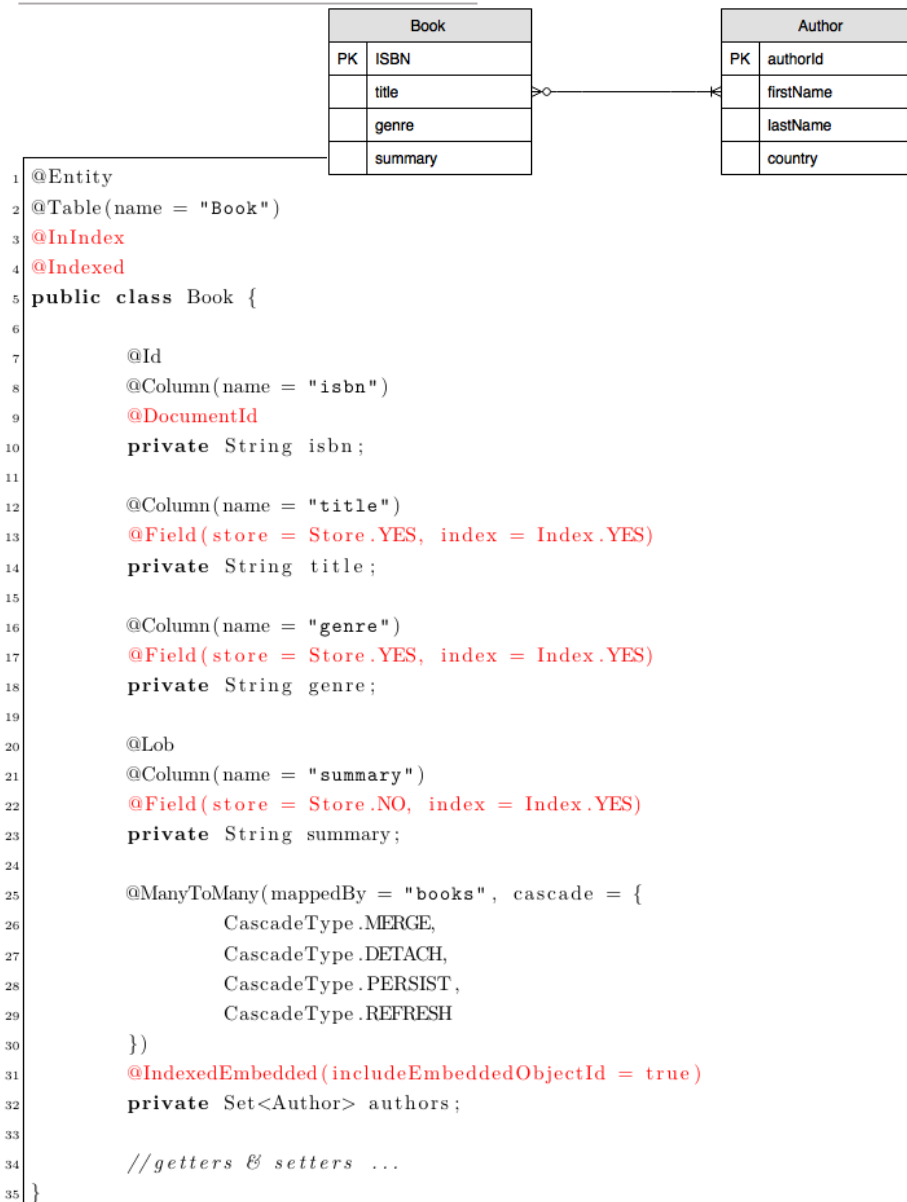
1 @Entity
2 @Table(name = "Book")
3
4
5 public class Book {
6
7     @Id
8     @Column(name = "isbn")
9
10    private String isbn;
11
12    @Column(name = "title")
13
14    private String title;
15
16    @Column(name = "genre")
17
18    private String genre;
19
20    @Lob
21    @Column(name = "summary")
22
23    private String summary;
24
25    @ManyToMany(mappedBy = "books", cascade = {
26        CascadeType.MERGE,
27        CascadeType.DETACH,
28        CascadeType.PERSIST,
29        CascadeType.REFRESH
30    })
31
32    private Set<Author> authors;
33
34    //getters & setters ...
35 }
  
```

```

1 @Entity
2 @Table(name = "Author")
3
4 public class Author {
5
6     @Id
7     @GeneratedValue(strategy = GenerationType.AUTO)
8     @Column(name = "authorId")
9
10    private Long authorId;
11
12    @Column(name = "firstName")
13
14    private String firstName;
15
16    @Column(name = "lastName")
17
18    private String lastName;
19
20    @Column(name = "country")
21
22    private String country;
23
24    @ManyToMany(cascade = {
25        CascadeType.MERGE,
26        CascadeType.DETACH,
27        CascadeType.PERSIST,
28        CascadeType.REFRESH
29    })
30    @JoinTable(name = "Author_Book",
31        joinColumns =
32            @JoinColumn(name = "authorFk",
33                referencedColumnName = "authorId"),
34        inverseJoinColumns =
35            @JoinColumn(name = "bookFk",
36                referencedColumnName = "isbn"))
37
38    private Set<Book> books;
39
40    //getters & setters ...
41 }
  
```





3. Standalone Version of Hibernate Search



```

1 @Entity
2 @Table(name = "Author")
3 @InIndex
4 public class Author {
5
6     @Id
7     @GeneratedValue(strategy = GenerationType.AUTO)
8     @Column(name = "authorId")
9     @DocumentId
10    private Long authorId;
11
12    @Column(name = "firstName")
13    @Field(store = Store.YES, index = Index.YES)
14    private String firstName;
15
16    @Column(name = "lastName")
17    @Field(store = Store.YES, index = Index.YES)
18    private String lastName;
19
20    @Column(name = "country")
21    @Field(store = Store.YES, index = Index.YES)
22    private String country;
23
24    @ManyToMany(cascade = {
25        CascadeType.MERGE,
26        CascadeType.DETACH,
27        CascadeType.PERSIST,
28        CascadeType.REFRESH
29    })
30    @JoinTable(name = "Author_Book",
31        joinColumns =
32            @JoinColumn(name = "authorFk",
33                referencedColumnName = "authorId"),
34        inverseJoinColumns =
35            @JoinColumn(name = "bookFk",
36                referencedColumnName = "isbn"))
37    @ContainedIn
38    private Set<Book> books;
39
40    //getters & setters ...
41 }

```

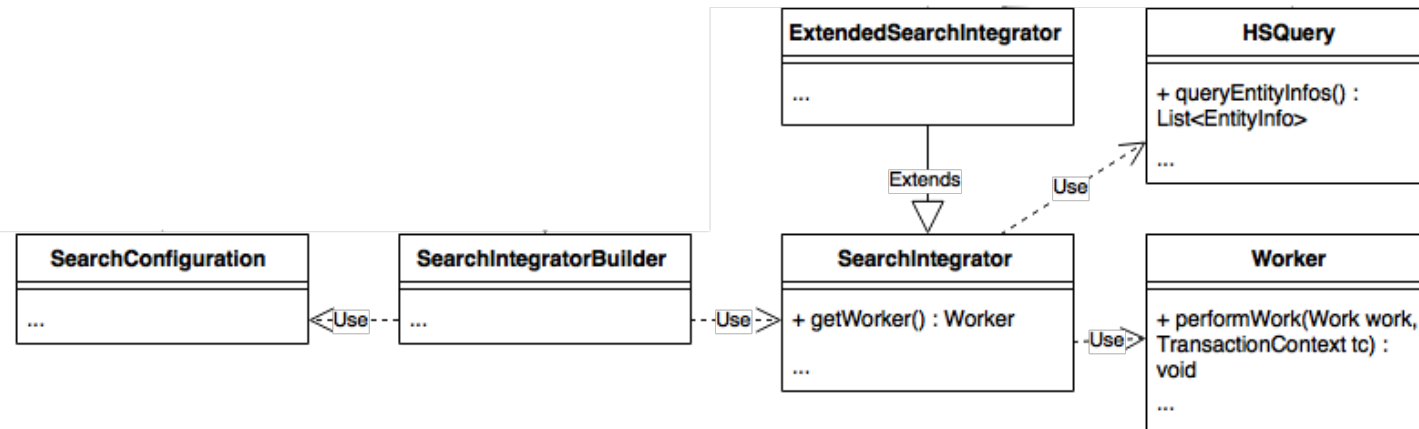


THURSDAY, 24 SEPTEMBER 2015

CONTACT

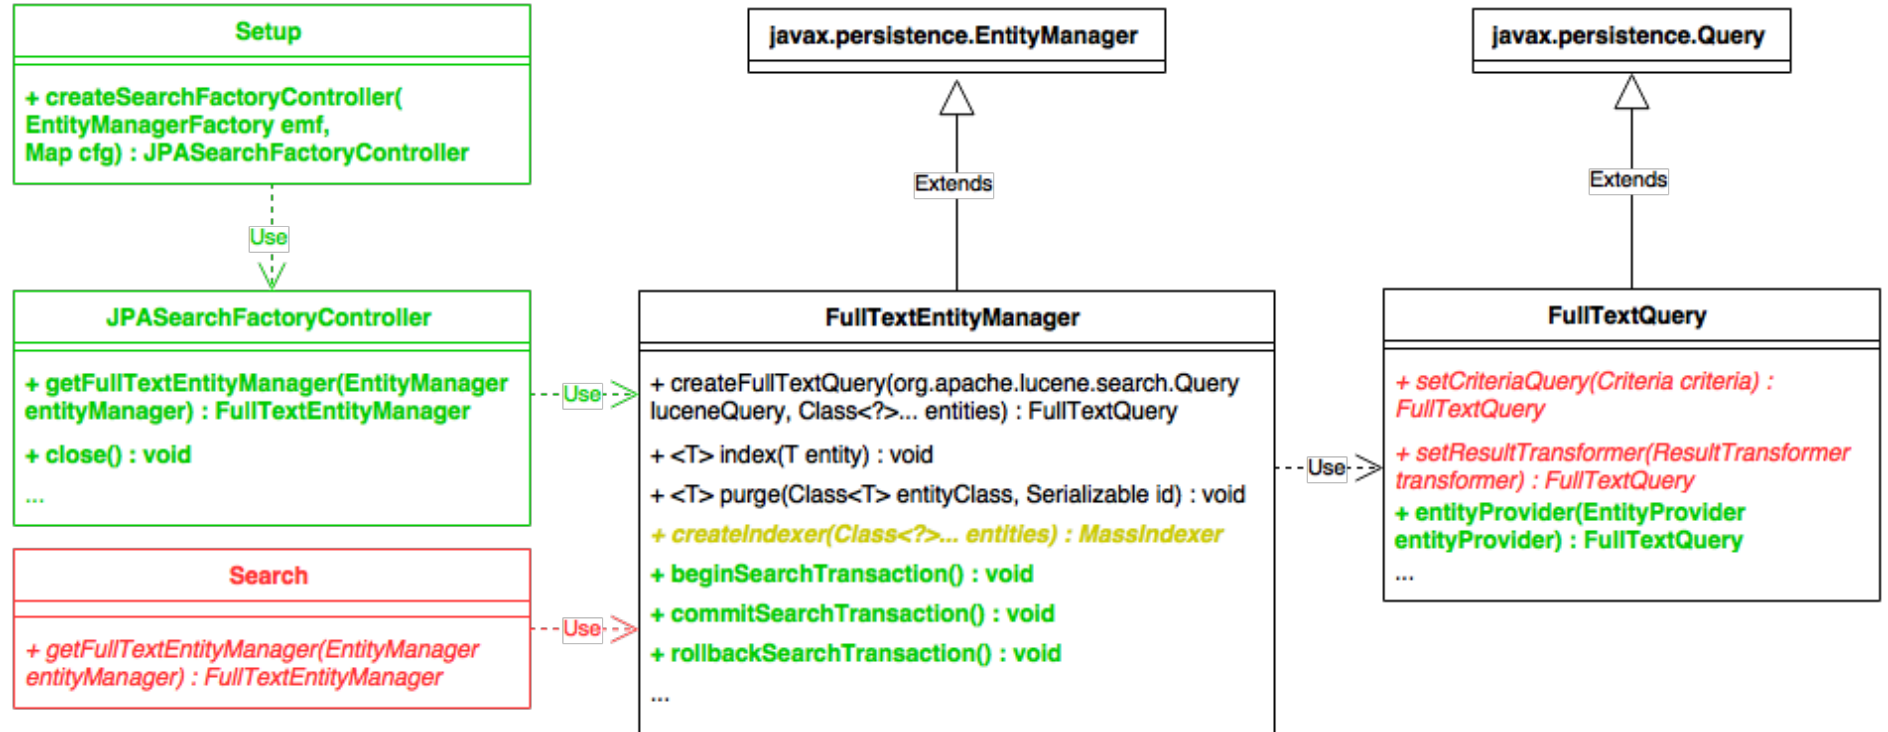
Martin Braun
E-Mail martinbraun123@aol.com

3. Standalone Version of Hibernate Search





4. JPA integration of the standalone version





THURSDAY, 24 SEPTEMBER 2015

CONTACT

Martin Braun
E-Mail martinbraun123@aol.com

5. Automatic Index Updating

- Index must be kept up-to-date when database changes
- Manual index updating is hard to maintain
- Hibernate Search ORM:
 - integrates with Hibernate ORM's event model
 - does not work with direct database changes (SQL on DB, custom queries)
- Candidates for Hibernate Search GenericJPA:
 - 1. JPA event based model synchronous
 - 2. Native event based model similar to Hibernate Search ORM synchronous
 - 3. Trigger based approach asynchronous



5. Automatic Index Updating - JPA events

```

1 public class EntityListener {
2
3     @PostPersist
4     public void persist(Object entity) {
5         //handle the event
6     }
7
8     @PostUpdate
9     public void update(Object entity) {
10        //handle the event
11    }
12
13    @PostDelete
14    public void delete(Object entity) {
15        //handle the event
16    }
17
18 }
```

```

1 @EntityListeners( { EntityListener.class } )
2 public class Book {
3
4     //...
5
6 }
```

Problems:

- JPA providers handle events differently



5. Automatic Index Updating - JPA events

```

1 EntityManager em = ...;
2
3 em.getTransaction().begin();
4
5 Book book = em.find( Book.class, "someIsbn" );
6 book.setTitle( "someNewTitle" );
7
8 // flushes, so we retrieve the Book with the changes from above
9 // => event is triggered
10 List<Book> allBooks =
11     em.createQuery( "SELECT b FROM Book b" ).getResultList();
12
13 // we have no way to get this event to revert the wrong index change
14 em.getTransaction().rollback();
    
```

Problems:

- JPA providers handle events differently
- Events are triggered on flush

→ unusable for index updating

THURSDAY, 24 SEPTEMBER 2015

CONTACT

Martin Braun
E-Mail martinbraun123@aol.com





5. Automatic Index Updating - Native Events

- All big JPA implementations (Hibernate ORM, EclipseLink, OpenJPA) have a native Listener mechanism
- Native update mechanisms in Hibernate Search GenericJPA:
 - Hibernate ORM
 - EclipseLink
- Same behaviour as Hibernate Search ORM's mechanism
 - do not work with direct database changes (SQL on DB, custom queries)
- Implementation is straight-forward, not part of thesis



THURSDAY, 24 SEPTEMBER 2015

CONTACT

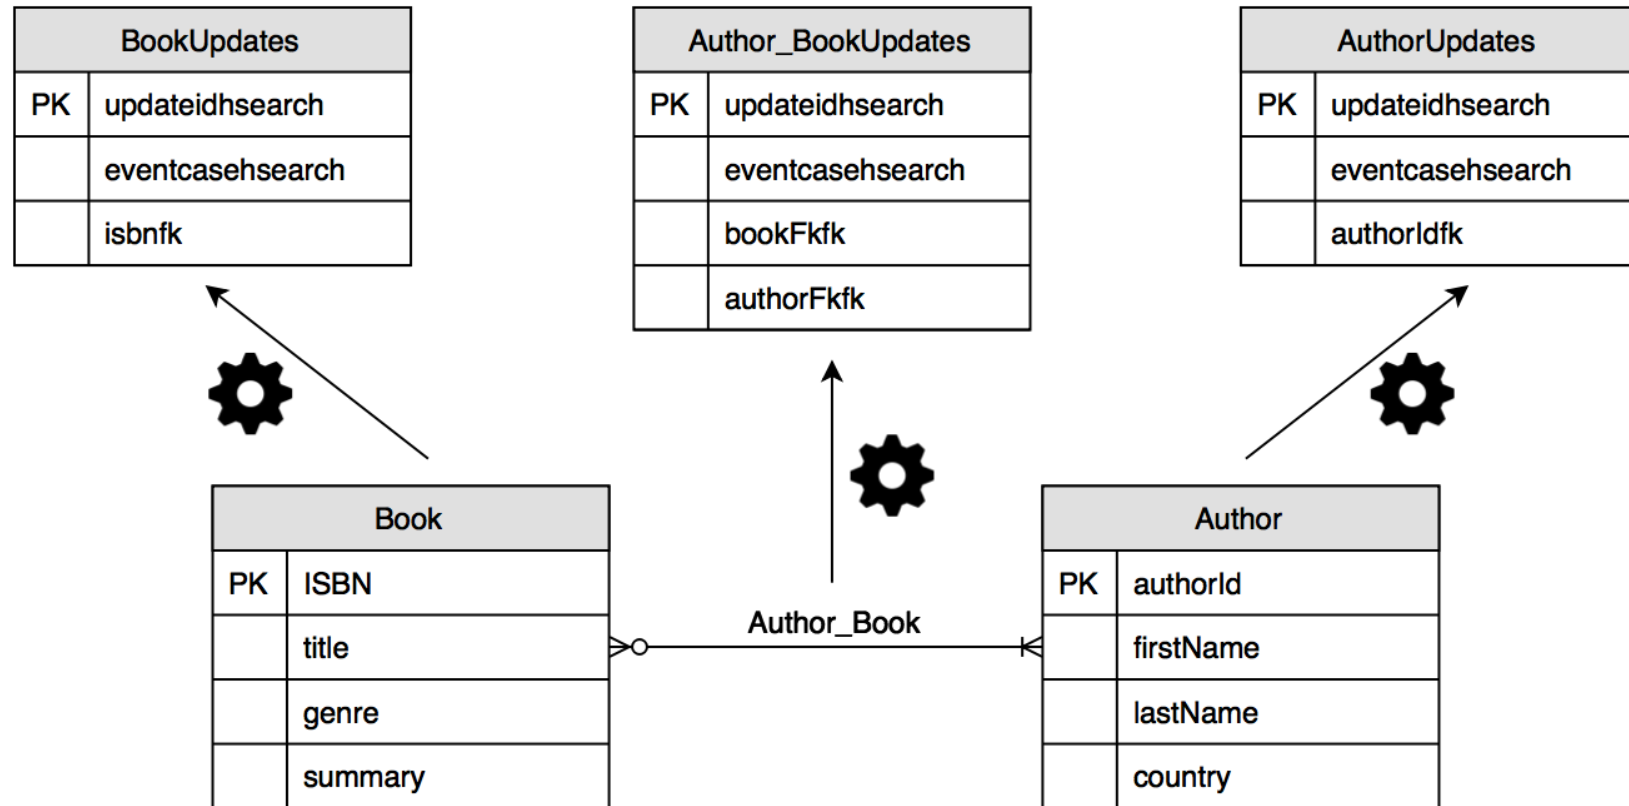
Martin Braun
E-Mail martinbraun123@aol.com

5. Automatic Index Updating - Triggers

- Triggers are supported by most RDBMSs
- Not standardized, code abstraction needed for generation
- Listen for changes directly in the database, write into auxiliary tables
 - changes with native SQL are recognized
 - support for legacy applications



5. Automatic Index Updating - Triggers





5. Automatic Index Updating - Triggers

```

1 @Entity
2 @InIndex
3 @Table(name = "Book")
4 @Indexed
5 @UpdateInfo(
6     tableName = "Book",
7     idInfos = @IdInfo(
8         columns = @IdColumn(
9             column = "isbn",
10            columnType = ColumnType.STRING
11        )
12    )
13 )
14 public class Book {
15
16     // ... unchanged.
17
18     // mapping table events handled on Author side
19
20     // getters & setters ...
21 }

```

```

1 @Entity
2 @InIndex
3 @Table(name = "Author")
4 @UpdateInfo(
5     tableName = "Author",
6     idInfos = @IdInfo(
7         columns = @IdColumn(
8             column = "authorId",
9             columnType = ColumnType.LONG
10        )
11    )
12 )
13 public class Author {
14
15     // ... unchanged.
16
17     @UpdateInfo(tableName = "Author_Book",
18         idInfos = {
19             @IdInfo(entity = Author.class,
20                 columns = @IdColumn(
21                     column = "authorFk",
22                     columnType = ColumnType.LONG
23                 )
24             ),
25             @IdInfo(entity = Book.class,
26                 columns = @IdColumn(
27                     column = "bookFk",
28                     columnType = ColumnType.STRING
29                 )
30             )
31         })
32     private Set<Book> books;
33
34     //getters & setters ...
35 }

```

THURSDAY, 24 SEPTEMBER 2015

CONTACT

Martin Braun
E-Mail martinbraun123@aol.com



CHAIR FOR DATABASES AND
INFORMATION SYSTEMS

PROF. DR.-ING. STEFAN JABLONSKI
DR. BERNHARD VOLZ

THURSDAY, 24 SEPTEMBER 2015

CONTACT

Martin Braun
E-Mail martinbraun123@aol.com

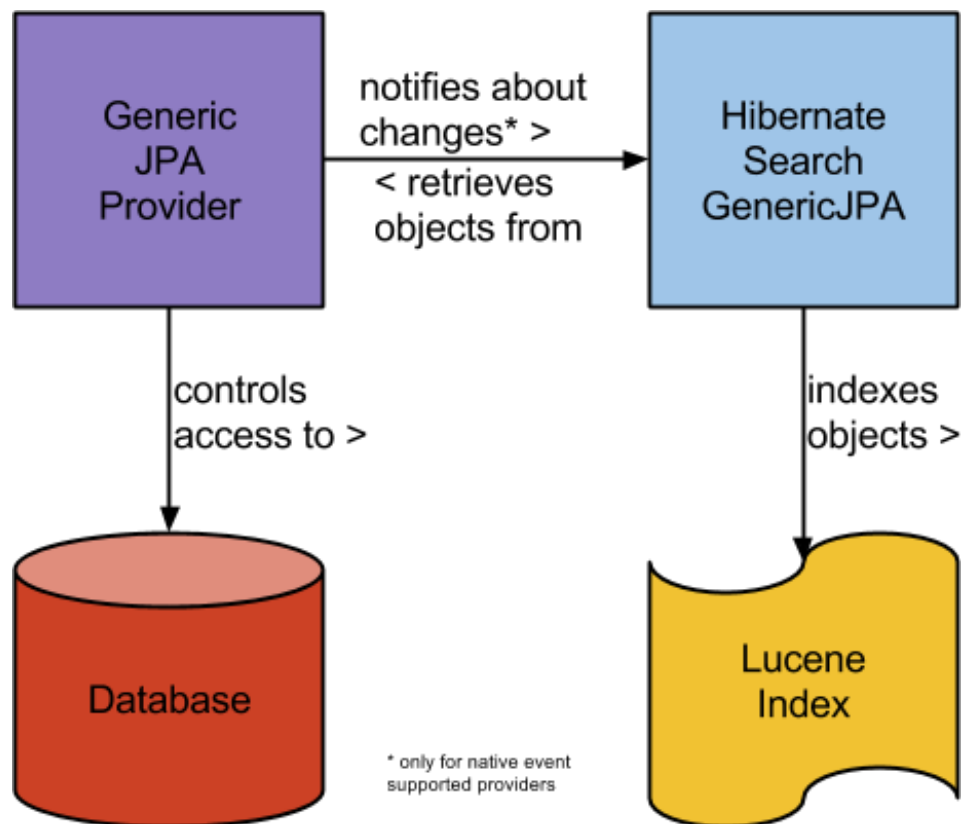
 DATENBANKEN UND
INFORMATIONSSYSTEME



6. Usage Example

7. Current situation

Hibernate Search without Hibernate ORM:



8. Outlook

- Stable proof of concept
- Source-code on GitHub:
<https://github.com/Hotware/Hibernate-Search-GenericJPA>
- Improvements in the trigger updating mechanism
- Merge with core Hibernate Search in November 2015





CHAIR FOR DATABASES AND
INFORMATION SYSTEMS

PROF. DR.-ING. STEFAN JABLONSKI
DR. BERNHARD VOLZ

THURSDAY, 24 SEPTEMBER 2015

CONTACT

Martin Braun
E-Mail martinbraun123@aol.com

 DATENBANKEN UND
INFORMATIONSSYSTEME



Lucene Basics

Documents		
id	field1	field2
1	fulltext search lucene	search
2	lucene search	java
3	fulltext java	fulltext lucene

Inverted Index		
Term		Occurences
Field	Value	
field1	fulltext	1,3
field1	search	1,2
field1	lucene	1,2
field1	java	3
field2	search	1
field2	java	2
field2	fulltext	3
field2	lucene	3